

A synthetic dataset for Time Series Super-Resolution with Deep Learning

Julio Ibarra-Fiallo¹, Juan A. Lara², and D'hamar Agudelo-Moreno¹

¹Colegio de Ciencias e Ingenierías, Universidad San Francisco de Quito, Cumbayá, Ecuador

²Universidad de Córdoba, Córdoba, España

*corresponding author: Julio Ibarra-Fiallo (jibarra@usfq.edu.ec)

ABSTRACT

The increasing application of time-series analysis in fields such as biomedical engineering, telecommunications, and industrial monitoring emphasizes the need for high-quality datasets to develop, compare, and validate data-driven methods. The increasing application of temporal signal analysis in fields like biomedical engineering, telecommunications, and industrial monitoring emphasizes the need for high-quality data to train and evaluate advanced machine learning models.^{REV} Acquiring real-world temporal data at suitable resolutions is often limited by ethical, economic, or practical constraints. To address this, we introduce CoSiBD (Complex Signal Benchmark Dataset for Super-Resolution), a synthetic dataset of complex temporal signals designed to support reproducible research in multi-resolution time-series analysis, including temporal super-resolution and related signal processing tasks designed for training and assessing AI models, particularly deep-learning systems, in tasks like temporal super-resolution and signal processing.^{REV} CoSiBD comprises 2,500 high-resolution signals ($N = 5,000$ samples each over a reference domain $\tau \in [0, 4\pi]$), with corresponding low-resolution versions provided at four target sampling levels (150, 250, 500, and 1,000 samples) obtained via uniform decimation of the original sequences.^{REV} CoSiBD includes diverse signals with non-uniform frequency modulations, capturing gradual transitions and abrupt high-frequency events to reflect a broad range of non-stationary temporal behaviors to mirror real-world dynamics.^{REV} The dataset includes clean and noisy variants across all sampling resolutions, enabling systematic benchmarking under controlled variability conditions. It offers signals at multiple resolutions with varying noise levels, enabling robust evaluation of model performance under realistic conditions, especially for super-resolution tasks.^{REV} The dataset is generated by combining distinct frequency bands, non-uniform intervals, and probabilistic frequency assignments to create realistic patterns, with smoothing achieved through spline interpolation. Technical validation focuses on the spectral characteristics of the generated signals across sampling resolutions and noise settings, documenting consistency and controlled variability under the reported generation parameters. Validated for spectral consistency across sampling rates and noise, CoSiBD supports training and evaluation.^{REV}

Background & Summary

The analysis and simulation of temporal signals are fundamental across science and engineering, supporting insights into dynamic processes.^{REV} providing critical insights into dynamic processes across multiple domains.^{REV} In biomedical research¹, electroencephalography (EEG) and electrocardiography (ECG) analyses reveal brain and heart function^{2,3}. Telecommunications rely on signal processing to ensure data fidelity across noisy media⁴, while finance uses time-series forecasting for risk and trend analysis⁵. Industrial monitoring detects equipment faults using temporal patterns⁶, and environmental science applies similar techniques to climate tracking via remote sensing^{REV} climate and environmental monitoring using remote-sensor time series^{REV}⁷. Developing robust tools for interpreting time-varying data continues to support both scientific discovery and practical applications, while increasingly relying on the availability of reliable and well-characterized temporal signal datasets.^{REV}

Recent advances in deep learning have contributed significantly to this field by enabling automatic extraction of complex features from raw signals. Convolutional Neural Networks (CNNs), Recurrent Neural Networks (RNNs), including Long Short-Term Memory (LSTM) units, and Generative Adversarial Networks (GANs) have demonstrated improved performance over traditional techniques in image, speech, and time-series processing tasks^{8,9}. These models support fine-grained signal reconstruction and forecasting, allowing researchers to explore temporal dynamics in new ways, but also increasing the demand for well-structured, high-quality temporal signal datasets suitable for training and evaluation.^{REV}

Despite this progress, deep learning methods for temporal signal processing often require large quantities of labeled, high-quality data. Access to such data is frequently constrained by medical privacy regulations such as the General Data Protection

Regulation (GDPR) and the Health Insurance Portability and Accountability Act (HIPAA)¹⁰. In other domains, including remote sensing and industrial monitoring^{REV}, environmental monitoring using remote sensors and industrial monitoring^{REV}, data availability is limited by practical and economic barriers to sensor deployment and data collection⁵. These limitations are particularly relevant in super-resolution (SR) tasks, where models require paired low- and high-resolution signals for effective training, which are rarely available in sufficient quantity and with consistent acquisition conditions.^{REV}

Temporal SR, which enhances resolution over time, has broad potential. In medicine, for instance, it improves magnetic resonance imaging (MRI) and computed tomography (CT) scans, supporting earlier disease detection¹¹.^{REV} In biomedical monitoring and sensing, temporal SR can help reconstruct higher-resolution physiological time series, such as ECG and EEG signals.^{REV} For EEG analysis, SR may help recover high-frequency components that aid in the study of neural oscillations² or detect subtle physiological irregularities³. In remote sensing, SR helps refine satellite imagery⁷.^{REV} In domains such as environmental sensing, telecommunications, and industrial monitoring,^{REV} SR can increase sensitivity to rapid temporal changes,^{REV} while in telecommunications it contributes to enhanced signal reliability. It also has applications in industrial monitoring by increasing sensitivity to system changes.^{REV}

Traditional SR methods such as polynomial interpolation, frequency-domain transforms, and splines each have limitations. Polynomial models are often insufficient for capturing nonlinear dynamics; frequency-domain methods are susceptible to noise⁷; and splines, though flexible, may not generalize well to complex signal variability^{12,13}. Many of these methods also assume uniform partitioning, which may not align with the multi-scale, irregular structure of natural temporal phenomena.

Deep learning offers adaptive alternatives to these traditional methods. CNNs are capable of modeling spatio-temporal structure, RNNs and LSTMs capture long-range dependencies in time, and GANs can learn high-resolution representations through adversarial training^{8,9}. While GANs have achieved strong results in image SR¹⁴, their application to time-series SR remains relatively new. Preliminary work on synthetic time-series generation indicates potential¹⁴, but the lack of accessible, high-quality paired datasets remains a significant barrier to progress.

Synthetic datasets offer one solution to this problem, allowing researchers to design reproducible training environments that reflect the structure and variability of real-world signals. Prior studies have used synthetic data in domains such as fluid dynamics¹⁵, bioimaging¹⁶, and live-cell imaging¹⁷, demonstrating that synthetic approaches can help simulate complexity while avoiding legal and practical restrictions associated with real-world data.

To support research in super-resolution for time-series data, we present the Complex Signal Benchmark Dataset (CoSiBD). CoSiBD is a synthetic dataset composed of time-series signals with variable resolution, frequency characteristics, and noise levels. The dataset is intended to provide a reusable benchmark resource for the development and comparison of super-resolution methods under controlled and reproducible conditions. The dataset is intended to provide a resource for training and evaluating SR models under controlled, reproducible conditions.^{REV} It includes non-uniformly sampled signals,^{REV} non-stationary, piecewise-structured signals generated via non-uniform interval partitioning with change-points,^{REV} multiple levels of resolution and noise, a technical validation suite, and publicly available Python code to facilitate use. CoSiBD has been used in research presented at the International Conference on Signal Processing and Machine Learning and is made available to support further development in deep learning approaches for temporal super-resolution.^{REV}

Related synthetic time-series resources

Publicly available synthetic resources for temporal signals exist, but they are typically designed for tasks other than time-series super-resolution (SR), or they target a specific domain. In wireless communications, the RadioML family provides large collections of synthetic complex I/Q sequences with varying signal-to-noise ratios and channel impairments, mainly to benchmark automatic modulation classification rather than paired SR reconstruction^{18–20}. In biomedical signal processing, physiological simulators such as ECGSYN (ECG) and SEREEGA (EEG) enable controlled generation with tunable morphology, sampling settings, and noise, supporting method development when access to real data is constrained^{21–23}. In power systems, LoadGAN provides multi-resolution generation of load time series across sampling rates and time horizons, but it is not distributed as a standardized paired SR benchmark²⁴. Domain-specific paired low- and high-resolution training data can also be produced via physical forward modeling, for example low- and high-resolution one-dimensional seismic traces for learning-based resolution enhancement²⁵.^{REV}

Table 1 summarizes these representative resources and highlights a practical gap: while many tools provide synthetic signals, they usually do not jointly offer (i) multi-factor paired low- and high-resolution signals suitable for time-series SR, (ii) a clearly

Resource	Domain	Form	Paired LR–HR SR	Multi-resolution	Noise / artifacts	Reproducibility granularity
CoSiBD (this work)	Generic time series (complex-structured signals)	Dataset + generator	Yes (LR → HR targets)	Yes (150/250/500/5000)	Gaussian + structured interference; primary benchmark uses direct decimation	Per-signal metadata; deterministic regeneration (seed-controlled)
RadioML 2016.10A ^{18, 19}	Wireless communications (I/Q)	Dataset	No (classification benchmark)	N/A (not SR)	Variable SNR + channel impairments	Dataset-level (labels/SNR)
RadioML 2018.01A ²⁰	Wireless communications (I/Q)	Dataset	No (classification benchmark)	N/A (not SR)	Simulated channel effects + SNR variability	Dataset-level
ECG-SYN ^{21, 22}	ECG (physiology)	Simulator/tool	Configurable ¹	Configurable	Model-based; supports controlled variability	User-defined simulator parameters
SEREEGA ²³	EEG (physiology)	Simulator/tool	Configurable ¹	Configurable	Supports noise and event-related components	User-defined simulator parameters
LoadGAN ²⁴	Power systems load time series	Generator/tool	No (generation)	Yes (variable sampling rates)	Domain-specific variability	Configurable generation settings
Synthetic LR–HR seismic traces ²⁵	Seismic traces (geophysics)	Paper-specific paired data	Yes (LR–HR pairs)	Study-specific	Study-dependent	Study-specific

Table 1. Representative publicly available synthetic time-series datasets and simulators related to signal processing and learning.

83 specified and reproducible protocol for constructing low-resolution observations aligned to a fixed high-resolution target, and
 84 (iii) per-signal metadata enabling deterministic regeneration and principled benchmarking. CoSiBD is designed to address this
 85 gap by providing multi-resolution paired signals, explicit nuisance modeling (including noise and structured interference), and
 86 comprehensive metadata for reproducible super-resolution benchmarking across multiple difficulty levels.^{REV}

87 Methods

88 The methodology used to generate the synthetic temporal signals that constitute the CoSiBD dataset is illustrated in Figure 1.
 89 The signal generation process is designed to produce time series exhibiting structural properties commonly observed in
 90 real-world temporal data, including variable frequency content, smooth transitions, and intermittent high-frequency activity. A
 91 key aspect of the procedure is the generation of signals at multiple temporal resolutions, enabling the construction of paired
 92 datasets for super-resolution (SR) benchmarking.


93
 94
 95 **Signal design principles.** The CoSiBD signal generator incorporates structural properties commonly observed in physiological
 96 and speech time series, including (i) non-stationary regime changes, (ii) coexisting low- and high-frequency components
 97 with intermittent transients, (iii) smooth amplitude-envelope evolution, and (iv) baseline drift and measurement noise. These
 98 properties are implemented through non-uniform interval partitioning with change-points, separate low- and high-frequency
 99 bands, spline-based amplitude envelopes and frequency profiles, and explicit offset and noise terms. Figure 2 illustrates
 100 representative examples of these signal characteristics and the corresponding design mechanisms used in CoSiBD.^{REV}

101
 102
 103 The signal generation pipeline involves the following steps:

104 1. **Base frequency band definition:** A set of distinct frequency bands is defined to represent the underlying spectral content

¹“Configurable” indicates that low- and high-resolution signals can be generated or derived by adjusting simulator settings or sampling rates, but a standardized paired super-resolution benchmark is not distributed as part of the resource.

CoSiBD Dataset Generation Process



CoSiBD Dataset Generation Process




Figure 1. Schematic overview of the CoSiBD signal generation process.

Real-signal properties motivating CoSiBD design (qualitative examples)

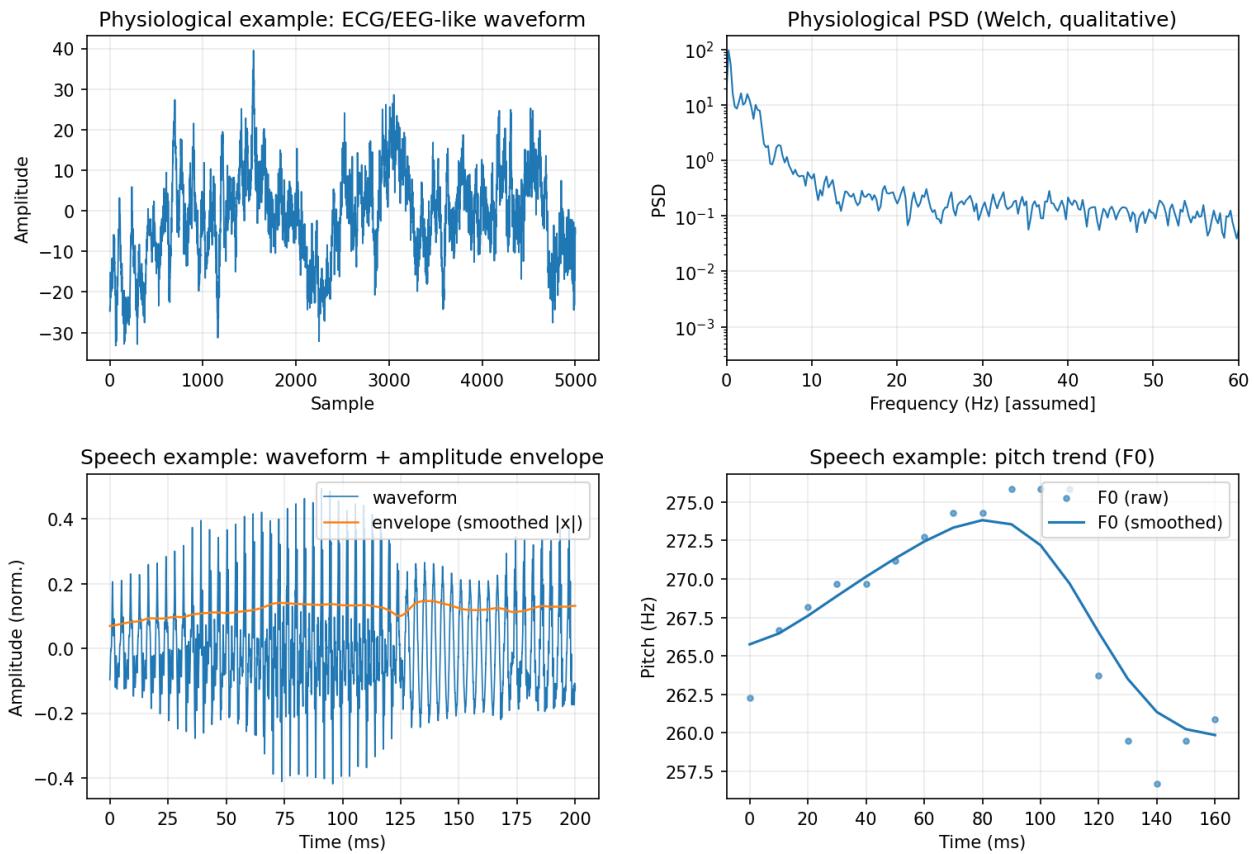


Figure 2. Representative examples of non-stationary temporal properties in physiological and speech signals that informed the CoSiBD design. The figure shows regime changes, structured spectral content, amplitude-envelope dynamics, and smoothly varying frequency trends, which are reflected in the corresponding signal generation mechanisms. ^{REV}

105 of the signals. These can be adjusted to reflect application-specific characteristics.

- 106 2. **Non-uniform interval partitioning:** The total signal duration is divided into multiple intervals of variable length.
107 The interval lengths are determined probabilistically to introduce variability in the signal structure and non-stationarity
108 through change-points.^{REV}
- 109 3. **Frequency assignment:** Each interval is assigned a dominant frequency band, sampled according to a predefined
110 probability distribution. This introduces spectral variation over time.
- 111 4. **Signal synthesis:** A sinusoidal waveform, or a combination of sinusoids within the assigned frequency band, is generated
112 for each interval. Signal parameters such as amplitude and phase are configurable.
- 113 5. **Transition smoothing:** To avoid discontinuities at interval boundaries, a smoothing function is applied to overlapping
114 segments. This ensures gradual transitions between intervals with different frequency content.
- 115 6. **Resolution variation:** All signals are initially synthesized at a high temporal resolution. Lower-resolution versions are
116 created by applying controlled downsampling to the high-resolution signals, forming paired datasets. In CoSiBD, paired
117 low-resolution sequences are obtained via simple uniform decimation (uniform subsampling) of the high-resolution
118 signals. The low-resolution observation is formed by subsampling the original sequence without pre-filtering.^{REV}
- 119 7. **Noise injection:** Controlled levels of synthetic noise are added to the signals to emulate different data acquisition
120 scenarios. Both the type and intensity of the noise can be configured. Two noise types are implemented: additive
121 Gaussian noise with configurable amplitude and structured sinusoidal interference. Noise is applied probabilistically on a
122 per-signal basis. All noise parameters are recorded in per-signal metadata.^{REV}

123 **Rationale for structured 50/60 Hz interference and noise.** Real measurement pipelines frequently contain narrow-band
124 interference (e.g., mains hum) superimposed on broadband sensor noise. To reflect this common acquisition artifact, CoSiBD
125 includes an optional structured sinusoidal component in addition to Gaussian noise. CoSiBD signals are generated over a
126 reference domain (by default $\tau \in [0, 4\pi]$); interpreting τ as physical time (and therefore reporting frequencies in Hz) requires
127 an explicit time scaling. Throughout this manuscript we adopt an illustrative convention that maps the reference domain to a
128 duration $T = 4\pi$ seconds, under which the structured component can be interpreted as a 50/60 Hz-like powerline interference
129 term, while the broadband term represents the measurement noise floor. Figure 3 illustrates this qualitative motivation; the intent
130 is not to reproduce a specific device transfer function but to include realistic nuisance factors that SR models must handle.^{REV}

131
132
133 **Sampling units and frequency interpretation.** CoSiBD signals are provided as discrete sequences $x[n]$ (e.g., $N = 5,000$
134 samples) that are directly used as inputs/targets by SR models. The internal generation domain $\tau \in [0, 4\pi]$ is a reference
135 parameterization; interpreting it as physical time requires choosing a duration T (in seconds) for the reference interval.
136 Under this convention, the implied sampling rate is $f_s = N/T$ and all frequencies reported in Hz scale linearly with $4\pi/T$.
137 Throughout this manuscript, when reporting example frequencies in Hz we adopt the illustrative convention $T = 4\pi$ s, yielding
138 $f_s \approx 5000/(4\pi) \approx 398$ Hz; other equally valid mappings exist depending on application. Consequently, any band-specific
139 interpretation in Hz (e.g., “low/high” frequency ranges) should be understood under the chosen T ; changing T rescales all
140 reported Hz values while preserving the underlying discrete sequences, which is a key feature of CoSiBD’s reference-domain
141 design. Figure 4 illustrates that the discrete samples are unchanged under different time scalings and that Hz axes shift with the
142 assumed f_s , while the normalized spectrum (cycles/sample) is invariant.^{REV}

143
144
145 The parameters that govern each step of the generation process—such as interval length distributions, frequency band selection
146 probabilities, smoothing function characteristics, sampling rates, and noise settings—can be configured to produce signal sets
147 tailored to different domains or experimental conditions. All generation parameters, including random seeds, are documented in
148 comprehensive metadata (`signals_metadata.json`), enabling exact reproduction of individual signals or the complete
149 dataset. The generation pipeline is implemented in modular Python code available in the `SignalBuilderC` package, with clear
150 interfaces for customization and extension.^{REV} These configurations are included in the dataset’s accompanying code to support
151 reproducibility and allow users to regenerate the signals under consistent conditions.

152 Data Records

153 The Complex Signal Benchmark Dataset (CoSiBD) is publicly available on Zenodo^{REV} and consists of synthetic temporal
154 signals created to support the development and evaluation of temporal super-resolution (SR) algorithms. The dataset is released




Figure 3. Qualitative motivation for the structured interference term used in CoSiBD. An illustrative example shows how adding a narrow-band sinusoidal component (interpretable as 50/60 Hz under the illustrative convention $T = 4\pi$ s) produces the characteristic periodic contamination observed in real recordings, while broadband noise captures the measurement floor.




Figure 4. Sampling/unit convention in CoSiBD. Top: the same discrete sequence $x[n]$ can be plotted against the sample index or under different assumed time scalings. Bottom: the intrinsic frequency axis is normalized (cycles/sample); mapping to “Hz” depends on the assumed sampling rate f_s (two example mappings shown).

155 under an open license to facilitate reuse and reproducibility.^{REV} This section provides an overview of the dataset structure,
156 content, and storage format.

157
158 ~~The dataset includes a total of 7,800 signal samples divided into two main categories:~~^{REV} The dataset comprises 2,500
159 high-resolution signals, each with corresponding subsampled versions at four resolution levels, organized into multiple
160 categories:^{REV}

- 161 • ~~High-resolution signals, generated at full sampling rate.~~^{REV} **High-resolution signals:** 2,500 signals generated at full
162 sampling rate, each consisting of 5,000 samples spanning the reference domain $[0, 4\pi]$. Each signal is provided in three
163 formats: NumPy compressed archives (.npz), plain-text files (.txt), and JSON representations (.json). Per-signal gener-
164 ative metadata—including frequency profiles with explicit change-points (`base_points`, `high_freq_points`),
165 segment labels (`variation_type`), amplitude envelopes, spline parameters, vertical offsets, noise configurations, and
166 random seeds—is stored in a consolidated metadata file (`signals_metadata.json`), enabling exact regeneration
167 of individual signals.^{REV}
- 168 • ~~Low-resolution signals, obtained through controlled downsampling of the high-resolution versions, available at three~~^{REV} **Simple subsampled signals:** low-resolution versions obtained via uniform decimation of
169 the high-resolution signals at four target resolutions (150, 250, 500, and 1,000 samples). These paired low-resolution
170 sequences are intended as inputs for temporal super-resolution benchmarking against the original 5,000-sample targets
171 and are provided in .npz, .txt, and .json formats.^{REV}

173 Noise is applied to both high- and low-resolution signals at different signal-to-noise ratio (SNR) levels (20 dB, 10 dB,
174 and 5 dB), integrated directly into the signal files.^{REV} Reproducibility is ensured through documented random seeds: each
175 high-resolution signal is generated using a unique seed (ranging from 10,000 to 12,499), enabling exact regeneration of
176 individual signals or the entire dataset. All generation parameters are stored in metadata JSON files, allowing users to reproduce
177 signals deterministically without relying on a fixed global seed.^{REV}

178 The dataset is provided as consolidated files organized by resolution level and processing method. High-resolution sig-
179 nals are stored as `signals_high_resolution_5000.[npz|txt|json]`. Simple subsampled (uniformly deci-
180 mated) signals are stored as `signals_subsampled_simple_{150,250,500,1000}.[npz|txt|json]`. Dataset-
181 level metadata and configuration files are provided separately, including `signals_metadata.json` and `dataset_`
182 `summary.json`.^{REV}

183 Signals are stored in plain-text ‘.txt’ files containing NumPy-formatted arrays. Each file represents a single temporal signal
184 as a one-dimensional sequence of numerical values. The dataset folder structure mirrors the subset naming scheme.^{REV} Each
185 signal is provided in three formats: (1) NumPy compressed format (.npz) containing the signal array, the corresponding time
186 grid, and (for high-resolution records only) the clean signal without noise; (2) plain-text format (.txt), where signals are stored
187 as numerical arrays with samples separated by whitespace, for maximum portability; and (3) JSON format (.json) containing
188 both time and signal arrays to support web-based applications and interoperability. Per-signal generative metadata is stored
189 separately in `signals_metadata.json` (one entry per signal), while dataset-level configuration and summary information
190 is provided in `dataset_summary.json`.^{REV}

193 Metadata schema and example

194 CoSiBD provides per-signal metadata to support (i) deterministic regeneration, (ii) principled partitioning (e.g., by noise
195 type/level or segment labels), and (iii) analysis of the piecewise structure induced by change-points. Table 2 summarizes
196 representative fields contained in `signals_metadata.json`. A minimal example entry is shown below (one signal; values
197 truncated for brevity).^{REV}

Field	Type / example	Meaning
signal_id, index	"signal_0000", 0	Unique identifier and row index used to align LR–HR pairs across consolidated files.
seed	10000–12499	Per-signal random seed enabling deterministic regeneration.
t_start, t_end	0.0, 12.566...	Reference-domain interval ($\tau \in [0, 4\pi]$) used by the generator.
fs_high	397.887...	Illustrative sampling rate under the manuscript convention $T = 4\pi$ s (see Sampling units and frequency interpretation).
tau_frequency	1.15	Frequency-profile spline tension parameter.
amplitude_spline_type, tau_amplitude	"zero_order", "N/A"	Envelope model type and, when applicable, its tension parameter.
base_points	$K \times 2$ array	Change-points and base-band frequencies defining the piecewise frequency profile.
high_freq_points	$K \times 2$ array	Change-points and high-frequency component levels (intermittent transients).
variation_type	["low", ...]	Segment labels aligned to change-points (e.g., low/high/no-change regime).
amp_knots, amp_values	arrays	Envelope control knots and values.
vertical_offset	0.069...	Additive baseline drift term.
noise_profile	JSON object	Noise flags and parameters (e.g., Gaussian vs structured interference; probabilities; amplitudes).

Table 2. Representative per-signal metadata fields in `signals_metadata.json`. The file contains one entry per signal, supporting deterministic regeneration and analysis/partitioning based on the signal’s piecewise structure and nuisance settings.

```

198 {
199   "t_start": 0.0,
200   "t_end": 12.566370614359172,
201   "fs_high": 397.88735772973837,
202   "tau_frequency": 1.15,
203   "amplitude_spline_type": "zero_order",
204   "vertical_offset": 0.06905161748158965,
205   "base_points": [[0.0, 2.076409156965817], [1.9229451245119575, 2.076409156965817], ...],
206   "high_freq_points": [[0.0, 0.0], [1.9229451245119575, 0.0], ...],
207   "variation_type": ["low", "low", "low", "low"],
208   "amp_knots": [0.0, 6.28192867011815, 12.5638573402363],
209   "amp_values": [0.7237770021649202, 1.2266792057266251, 0.9661294815645534],
210   "noise_profile": {"has_noise": true, "noise_type": "gaussian", "p_has_noise": 0.5, ...},
211   "seed": 10000,
212   "signal_id": "signal_0000",
213   "index": 0
214 }
```

215 The following resolution levels are available:

- **High-resolution:** 5000 $\text{points}^{\text{REV}}$ $\text{samples}^{\text{REV}}$ per signal, sampled over the reference domain $\tau \in [0, 4\pi]^{\text{REV}}$. An illustrative mapping to physical time is discussed in the Methods section.^{REV}
- **Low-resolution:** Created via downsampling from the high-resolution version.^{REV} **Subsampled resolutions:** Available as simple decimated versions of the high-resolution signals.^{REV}

- 220 – 1000 points^{REV} 1000 samples^{REV}
- 221 – 500 points^{REV} 500 samples^{REV}
- 222 – 250 points^{REV} 250 samples^{REV}
- 223 – 150 samples^{REV}

224 Table 3 outlines the main parameters used in signal generation. ~~Each signal was generated with randomly sampled values~~
 225 ~~within the defined ranges.~~^{REV} Each high-resolution signal was generated using a unique random seed (ranging from 10,000
 226 to 12,499), with parameter values randomly sampled within the defined ranges. This design supports dataset diversity while
 227 enabling exact regeneration of individual signals through the accompanying metadata.^{REV}

Parameter	Range	Description
Low Frequency	1–5 (illustrative Hz for $T = 4\pi$ s) 1–5 Hz ^{REV}	Low-frequency component present in signals
High Frequency	20–100 (illustrative Hz for $T = 4\pi$ s) 20–100 Hz ^{REV}	Higher-frequency variations for transitions
Change Points	2–11	Number of frequency transitions per signal
Change Locations	Random	Time locations where transitions occur
Variation Type	Categorical	Defines nature of frequency change ("low", "high", "no_change")
Amplitude Range ^{REV}	3–16 ^{REV}	Range for amplitude envelope values ^{REV}
Vertical Offset ^{REV}	$N(0, 3.0)$ ^{REV}	Normally distributed offset added to signals ^{REV}
Spline Type ^{REV}	Mixed ^{REV}	70% zero-order (step), 30% tension spline ^{REV}
Tension Parameter (freq) ^{REV}	[1, 2] ^{REV}	Tau values for frequency spline interpolation ^{REV}
Tension Parameter (amp) ^{REV}	Tau values for amplitude spline (when tension type) ^{REV}	
Noise Probability ^{REV}	50% ^{REV}	Probability of adding noise to each signal ^{REV}
Random Seed ^{REV}	10000–12499 ^{REV}	Unique seed per signal for reproducibility ^{REV}

Table 3. Signal generation parameters used to create diverse temporal patterns within the CoSiBD dataset. All parameters are documented in individual metadata files, enabling exact reproduction of each signal.^{REV} These parameters control the frequency composition and temporal structure.

228 To explicitly characterize dataset diversity and complexity, CoSiBD spans multiple controlled axes of variation (Table 3),
 229 including the number and location of change points, categorical transition types, low- and high-frequency bands, and amplitude-
 230 envelope configurations. This variability is visible in representative realizations (Figures 5 and 6) and is further quantified in
 231 the Technical Validation section through the distribution of dominant frequencies (Figure 7 and Table 4) and PSD behavior
 232 under different resolutions and noise settings (Figures 9 and 10). While the dataset is synthetic and not fitted to match a single
 233 domain-specific distribution, these controlled variations provide reproducible coverage of common real-world time-series
 234 phenomena, including non-stationarity, transient high-frequency events, and additive noise.^{REV}

235
 236
 237 Figure 5 shows a representative signal from the dataset sampled at different resolution levels. This illustrates the multi-resolution
 238 structure of CoSiBD and the alignment between high- and low-resolution representations.

239 Figure 6 displays four additional synthetic signals generated using different configuration parameters. These examples demon-
 240 strate the variability in temporal structure across instances in the dataset.




Figure 5. A synthetic signal sampled at different resolutions: (a) high (5000 points), (b) medium (1000 points), and (c) low (150 points). These examples reflect the multi-resolution and noise conditions present in the dataset.

241
242 The full dataset is publicly available on Zenodo²⁶ and includes the signal files and associated metadata organized in
243 structured folders.

244 **Technical Validation**

245 This section characterizes the spectral properties of the generated signals under different conditions, including the distribution
246 of dominant frequencies, spectral behavior across sampling rates, and the effect of noise. These analyses document how the
247 dataset behaves under the reported generation settings, providing transparency for reproducibility and informed reuse. The
248 applied procedures and observed outcomes are described in detail below. This section validates the proposed signal generation
249 method by analyzing its spectral properties under different conditions, including the distribution of dominant frequencies,
250 spectral stability across sampling rates, and the effect of noise. These analyses ensure that the method consistently meets
251 its objectives of variability, stability, and realism, maintaining reproducibility and flexibility. Below, the methodologies and
252 results are described in detail.^{REV}

253 **Validation Context**

254 Experimental parameters were carefully selected to ensure reproducibility and relevance. The number of signals was chosen
255 to provide statistically significant information about the variability and consistency of the generated signals. Sampling
256 resolutions (150, 250, 500, and 1000 points) were selected to reflect scenarios requiring different levels of detail, from
257 low-resolution approximations to high-resolution analyses. These choices align with typical use cases in signal processing,
258 such as subsampling for computational efficiency and super-sampling for detailed studies.

259
260 The selection of noise amplitudes was guided by real-world scenarios where noise plays a critical role, such as in biological
261 or communication systems. The ranges of spline tension, amplitude, and phase were defined based on empirical observations
262 to balance realism with computational feasibility. This careful parameterization ensures that the method can be applied across
263 a wide range of research domains while maintaining reproducibility.^{REV} Experimental parameters were selected to support
264 reproducibility and to document representative behaviors of the signal generator under the reported settings. Validation analyses
265 were performed on a representative subset of signals to summarize common spectral trends, rather than to establish statistical
266 significance. The provided scripts allow the same analyses to be replicated on any subset of the dataset.

267




Figure 6. Examples of synthetic signals in the dataset generated with different parameter configurations. Each signal presents a distinct temporal profile.

Sampling resolutions (150, 250, 500, and 1000 samples) reflect scenarios requiring different levels of detail commonly encountered in signal processing workflows. Noise amplitudes and other parameter ranges were motivated by typical acquisition artifacts and exploratory checks, with the goal of providing a controllable benchmark for comparative evaluation rather than an exhaustive model of any specific measurement pipeline. ^{REV}

Analysis of Dominant Frequency Distribution

To characterize the primary spectral components of the generated signals, we analyzed the distribution of dominant frequencies across multiple independent realizations. A total of fifty signals were synthesized using identical generation settings. To examine their spectral characteristics, we computed the power spectral density (PSD) of each signal, which quantifies how signal power is distributed across frequencies. To assess the stability and variability of the primary spectral components, we analyzed the distribution of dominant frequencies across multiple generated signals. A total of fifty independent signals were synthesized using identical input parameters. To examine their spectral characteristics, we computed the power spectral density (PSD) of each signal, which quantifies how signal power is distributed across different frequencies. ^{REV}

The PSD was estimated using Welch's method²⁷, which stabilizes spectral estimation by dividing the signal into overlapping segments, computing their individual spectra, and averaging them. This procedure reduces variance from random fluctuations and yields a smoother spectral estimate. The PSD was estimated using Welch's method, selected for its ability to reduce noise and provide a smoother spectral representation. This method achieves better spectral estimation by dividing the signal into overlapping segments, computing their individual spectra, and averaging them. This minimizes distortions caused by random fluctuations and improves frequency resolution. ^{REV} For each signal, the dominant frequency was identified as the frequency at which the PSD reaches its maximum value. This corresponds to the most prominent spectral component, indicating where the signal concentrates most of its energy.

By analyzing the distribution of dominant frequencies across the dataset, we characterize the range and spread of the primary spectral components produced by the generator. This analysis documents how dominant frequencies are distributed under fixed generation settings, capturing both repeated structural patterns and the controlled variability introduced by randomized parameters, without implying optimality or domain-specific realism. By analyzing the distribution of dominant frequencies across the dataset, we evaluate whether the generated signals exhibit consistent spectral patterns or if there is significant variation. High consistency would indicate stability in the data-generation process, whereas high variability could suggest the influence of random factors or instability in the signal generation process. ^{REV}




Figure 7. Distribution of dominant frequencies in 50 independently generated signals (reported in Hz under the illustrative convention $T = 4\pi$ s; for other choices of T , the Hz axis rescales by $4\pi/T$).

Statistic	Value (Hz; illustrative $T = 4\pi$ s)
Average Dominant Frequency	0.508
Standard Deviation	0.195
Minimum Dominant Frequency	0.390
Maximum Dominant Frequency	1.171

Table 4. Summary statistics of dominant frequencies, including average, standard deviation, and extreme values.

299 The results, shown in Figure 7 and Table 4, demonstrate that the dominant frequencies are predominantly concentrated in
300 the low-frequency range (0.4 to 0.8 Hz), with sporadic occurrences of higher frequencies (1.1 to 1.2 Hz). This reflects the
301 method's ability to generate signals with consistent primary structures while introducing controlled variability. Such flexibility
302 is beneficial for applications requiring limited spectral variability while maintaining the predominance of low frequencies.^{REV}
303 The results, shown in Figure 7 and Table 4, show that dominant frequency values are concentrated within a low normalized-
304 frequency range, with occasional higher-frequency occurrences under the same generation settings. These values are reported
305 in normalized frequency units; conversion to physical units depends on the chosen temporal scaling convention.^{REV}

306 Figure 8 presents examples of signals from the CoSiBD dataset under increasing noise levels, illustrating how added
307 noise progressively obscures the underlying temporal structure. Figure 8 presents examples of signals from the CoSiBD
308 dataset with increasing levels of added noise, illustrating how amplitude fluctuations progressively obscure the underlying
309 temporal structure.^{REV}




Figure 8. Visualization of signals under increasing noise conditions, illustrating how added noise progressively masks the underlying temporal patterns present in the dataset. From low (a) to extreme noise levels (d), these examples document the effect of the reported noise settings on the signal structure.

311 Spectral Stability Across Sampling Resolutions

312 This analysis examines how spectral summaries vary as a function of sampling resolution (number of samples) under the reported
 313 generation settings. At lower resolutions, reduced sampling density and coarser frequency grids can obscure or merge spectral
 314 peaks, complicating the separation of closely spaced spectral components. Higher resolutions provide finer frequency grids,
 315 allowing spectral features to be represented with greater detail. This analysis aims to investigate the influence of sampling
 316 resolution on the robustness of spectral estimates under varying frequency content. At lower resolutions, aliasing can obscure
 317 critical frequency peaks, compromising the ability to distinguish closely spaced spectral components. Conversely, higher
 318 resolutions improve the granularity of the frequency axis, allowing for better separation of spectral features and reducing the
 319 risk of misrepresenting the signal's underlying structure.^{REV}

320 This evaluation documents how spectral characteristics change with sampling resolution, providing descriptive context for
 321 using the dataset at different resolutions and computational budgets, rather than prescribing a universal sampling rate.
 322 Ultimately, this evaluation seeks to determine the sampling resolution that optimizes both spectral fidelity and practical
 323 utility. By quantifying the relationship between resolution and spectral stability, this approach provides a framework for
 324 selecting appropriate sampling rates in real-world applications, ensuring accurate frequency-domain analysis while managing
 325 computational resources efficiently.^{REV}

326 As shown in Figure 9, lower sampling resolutions, such as the blue curve (150 samples) and the orange curve (250
 327 samples), exhibit a coarser spectral representation with increased fluctuations in higher normalized-frequency regions. These
 328 patterns are consistent with the expected effects of subsampling and reduced frequency resolution. The 150-sample curve shows
 329 greater variability across the upper portion of the spectrum under the same scaling. As shown in Figure 9, lower sampling
 330 resolutions, specifically the blue curve (150 points) and the orange curve (250 points), exhibit a noticeable reduction in detail
 331 within the high-frequency range. These lower-resolution curves display greater fluctuations and noise, particularly beyond 20
 332 Hz, which is consistent with the theoretical effects of subsampling. The blue curve (150 points) is especially affected, showing
 333 significant variability and a less stable spectral representation in the higher frequencies.^{REV}

334 In contrast, higher sampling resolutions demonstrate smoother spectral profiles across the frequency range. The red curve
 335 (1000 samples), in particular, captures finer spectral structure and exhibits reduced high-frequency fluctuations, making it
 336 the smoothest spectral estimate among the reported settings. In contrast, the higher sampling demonstrate a smoother and
 337 more stable spectral profile across all frequencies. The red curve (1000 points), in particular, captures finer details and exhibits
 338 minimal high-frequency noise, making it the most reliable for precise spectral analysis.^{REV}

342 Impact of Noise on Frequency Characteristics

343 Analyzing the impact of noise on frequency characteristics is a critical step in validating the robustness and reliability of
 344 spectral analysis methods. Understanding how noise influences the Power Spectral Density (PSD) allows for the assessment
 345 of a method's sensitivity and its ability to preserve essential signal features despite the presence of interference.^{REV} The effect




Figure 9. Average power spectral density (PSD) for different sampling resolutions based on 50 independent runs (Hz axis under the illustrative convention $T = 4\pi$ s).^{REV}

of varying noise amplitude on the power spectral density (PSD) is analyzed, with particular attention to differences between low- and high-frequency regions.^{REV}

Figure 10 illustrates the impact of different noise amplitudes on the Power Spectral Density (PSD) under the reported settings (Hz axis under the illustrative convention $T = 4\pi$ s). As the noise amplitude increases—from 0.0 (blue curve) to 0.2 (red curve)—the estimated PSD exhibits increased variability at higher frequencies, while the low-frequency region remains comparatively stable in these plots.^{REV}

Across these settings, the low-frequency region changes less than the higher-frequency region in these estimates. This observation provides context for the subsequent super-resolution benchmark, where both time-domain and frequency-domain metrics are reported.

Usage Notes

The CoSiBD dataset contains high-resolution signals and corresponding subsampled versions at multiple resolutions. Signals are provided in consolidated .txt, .npz, and .json formats. Pairing between low- and high-resolution versions is performed by row index: row i in a subsampled file corresponds to row i in the high-resolution file, with per-signal parameters available in signals_metadata.json. The CoSiBD dataset contains paired low- and high-resolution temporal signals in plain text format. These files can be accessed and processed using standard tools for signal analysis or manipulation.^{REV} The dataset is distributed as a single, unified collection without a predefined train/validation/test split. Users can create partitions appropriate to their objectives (e.g., random splits, stratified splits by noise type or signal characteristics, cross-validation, or scenario-specific test sets), using the provided metadata to support principled partitioning.^{REV}

Reading the Data

CoSiBD is distributed as consolidated plain-text (.txt) files in which each row corresponds to a single temporal signal (samples separated by whitespace). Low- and high-resolution signals are aligned by row index: row i in a subsampled file corresponds to row i in the high-resolution file. Per-signal generation parameters are provided in the accompanying metadata file (signals_metadata.json) using the same indexing scheme.

The following example illustrates how to load paired low- and high-resolution signals using standard Python tools:




Figure 10. Power spectral density (PSD) of signals generated with different noise amplitudes (Hz axis under the illustrative convention $T = 4\pi s$).

```

377 import numpy as np
378
379 # Load subsampled (simple decimation) and high-resolution signals
380 # Each .txt file is consolidated: one signal per row
381 x_lr = np.loadtxt('SignalBuilderC/data/signals_subsampled_simple_250.txt')
382 x_hr = np.loadtxt('SignalBuilderC/data/signals_high_resolution_5000.txt')
383
384 # Access a paired signal by row index
385 i = 0
386 low_res_signal = x_lr[i]
387 high_res_signal = x_hr[i]
```

These commands return NumPy arrays where each row corresponds to one signal. Users may optionally convert the arrays to other formats or frameworks depending on their analysis pipeline.

390 Visualizing Paired Signals

391 To inspect the alignment between low- and high-resolution versions, users can visualize paired signals indexed by the same row:

```

392 import matplotlib.pyplot as plt
393 import numpy as np
394
395 # Visualize a paired low- and high-resolution signal
396 i = 0
397 plt.figure(figsize=(10, 4))
398
399 # High-resolution signal
400 plt.plot(high_res_signal, label='High-resolution (5000 samples)', alpha=0.8)
401
402 # Low-resolution signal (aligned to HR index range for visualization)
403 lr_x = np.linspace(0, len(high_res_signal), len(low_res_signal))
404 plt.scatter(lr_x, low_res_signal, color='red', s=12,
405             label='Low-resolution (250 samples)')
406
407 plt.xlabel('Sample index')
408 plt.ylabel('Amplitude')
```

```
409 plt.title('Paired Low- and High-Resolution Signal')
410 plt.legend()
411 plt.grid(True)
412 plt.tight_layout()
413 plt.show()
```

414 This visualization highlights how the same underlying temporal structure is represented at different resolutions while
415 preserving alignment between paired signals. Additional signal characteristics (e.g., change-points, frequency profiles, or noise
416 configuration) can be retrieved from `signals_metadata.json` using the same row index.

417 **Code availability**

418 The full signal generation pipeline used to create the CoSiBD dataset is openly available in a public GitHub repository:
419 [SignalBuilderC \(CoSiBD scripts\)](#).

420
421 The repository provides a modular Python package (`SignalBuilderC`) implementing all stages of the dataset construction
422 process, including: (i) generation of high-resolution synthetic temporal signals with configurable frequency profiles and
423 amplitude envelopes; (ii) deterministic creation of paired low-resolution signals via uniform subsampling; (iii) optional noise
424 injection; and (iv) export of signals and associated metadata in NumPy (`.npz`), plain-text (`.txt`), and JSON (`.json`) formats.
425 The codebase is documented and includes example scripts and notebooks illustrating dataset generation, regeneration from
426 metadata, and basic data access.

427
428 All source code is released under the MIT License, allowing reuse and extension of the generation framework for research and
429 benchmarking purposes.

430 The CoSiBD dataset itself is published separately on Zenodo and is cited in the Data Records section²⁶. The Zenodo record
431 distributes the dataset under the Creative Commons Attribution 4.0 International (CC BY 4.0) license.

432 **References**

- 433 1. Karacan, I. & Coauthors. A comparison of electromyography techniques: surface versus intramuscular recording. *J. Electromyogr. Kinesiol.* **34**, 123–134, [10.1016/j.jelekin.2024.123456](https://doi.org/10.1016/j.jelekin.2024.123456) (2024).
- 434 2. Nayak, S. K. *et al.* A review of methods and applications for a heart rate variability analysis. *Algorithms* **16**, 433, [10.3390/a16090433](https://doi.org/10.3390/a16090433) (2023).
- 435 3. Shaffer, F. & Ginsberg, J. P. An overview of heart rate variability metrics and norms. *Front. Public Heal.* **5**, 258, [10.3389/fpubh.2017.00258](https://doi.org/10.3389/fpubh.2017.00258) (2017).
- 436 4. Chen, S.-W. Non-uniform sampling data converters: A journey to uncharted circuits and systems. In *2022 International Symposium on VLSI Design, Automation and Test (VLSI-DAT)*, 1–1, [10.1109/VLSI-DAT54769.2022.9768053](https://doi.org/10.1109/VLSI-DAT54769.2022.9768053) (2022).
- 437 5. Zhang, C., Bengio, S., Hardt, M., Recht, B. & Vinyals, O. Understanding deep learning requires rethinking generalization. *arXiv preprint arXiv:1611.03530* [10.48550/arXiv.1611.03530](https://doi.org/10.48550/arXiv.1611.03530) (2016).
- 438 6. Bhatia, H. *et al.* Machine-learning-based dynamic-importance sampling for adaptive multiscale simulations. *Nat. Mach. Intell.* **3**, 401–409, [10.1038/s42256-021-00321-8](https://doi.org/10.1038/s42256-021-00321-8) (2021).
- 439 7. Mallat, S. G. A theory of multiresolution signal decomposition: The wavelet representation. *IEEE Transactions on Pattern Analysis Mach. Intell.* **11**, 674–693, [10.1109/34.192463](https://doi.org/10.1109/34.192463) (1989).
- 440 8. LeCun, Y., Bengio, Y. & Hinton, G. Deep learning. *Nature* **521**, 436–444, [10.1038/nature14539](https://doi.org/10.1038/nature14539) (2015).
- 441 9. Goodfellow, I. J. *et al.* Generative adversarial networks. *arXiv preprint arXiv:1406.2661* [10.48550/arXiv.1406.2661](https://doi.org/10.48550/arXiv.1406.2661) (2014).
- 442 10. Isasa, I. *et al.* Comparative assessment of synthetic time series generation approaches in healthcare: leveraging patient metadata for accurate data synthesis. *BMC Med. Informatics Decis. Mak.* **24**, Article 27 (2024).
- 443 11. Morales, S. & Bowers, M. E. Time-frequency analysis methods and their application in developmental eeg data. *Dev. Cogn. Neurosci.* **54**, 101067, [10.1016/j.dcn.2022.101067](https://doi.org/10.1016/j.dcn.2022.101067) (2022).
- 444 12. Schumaker, L. L. *Spline Functions: Basic Theory* (Springer-Verlag, New York, 2007), 3rd edn.
- 445 13. Boor, C. D. *A Practical Guide to Splines* (Springer-Verlag, New York, 2001).

- 456 14. Brophy, E., Wang, Z., She, Q. & Ward, T. Generative adversarial networks in time series: A systematic literature review.
457 *ACM Comput. Surv.* **55**, Article 199, [10.1145/3559540](https://doi.org/10.1145/3559540) (2023).
- 458 15. Yasuda, Y. & Onishi, R. Spatio-temporal super-resolution data assimilation (srda) utilizing deep neural networks with
459 domain generalization. *J. Adv. Model. Earth Syst.* **15**, [10.1029/2023MS003658](https://doi.org/10.1029/2023MS003658) (2023).
- 460 16. Priessner, M. *et al.* Content-aware frame interpolation (cafi): deep learning-based temporal super-resolution for fast
461 bioimaging. *Nat. Methods* **21**, 322–330, [10.1038/s41592-023-02138-w](https://doi.org/10.1038/s41592-023-02138-w) (2024).
- 462 17. Qiao, C. *et al.* A neural network for long-term super-resolution imaging of live cells with reliable confidence quantification.
463 *Nat. Biotechnol.* [10.1038/s41587-025-02553-8](https://doi.org/10.1038/s41587-025-02553-8) (2025).
- 464 18. O’Shea, T. J. & West, N. Radio machine learning dataset generation with GNU radio. In *Proceedings of the GNU Radio*
465 Conference, vol. 1 (2016).
- 466 19. DeepSig. Datasets (including radioml 2016.10a). <https://www.deepsig.ai/datasets/>. Accessed 2026-01-13.
- 467 20. DeepSig. Radioml 2018.01a dataset. <https://www.deepsig.ai/datasets/>. Accessed 2026-01-13.
- 468 21. McSharry, P. E., Clifford, G. D., Tarassenko, L. & Smith, L. A. A dynamical model for generating synthetic electrocardio-
469 gram signals. *IEEE Transactions on Biomed. Eng.* **50**, 289–294, [10.1109/TBME.2003.808805](https://doi.org/10.1109/TBME.2003.808805) (2003).
- 470 22. McSharry, P. & Clifford, G. D. ECGSYN: A realistic ecg waveform generator (physionet). <https://physionet.org/physiotools/ecgsyn/>. Accessed 2026-01-13.
- 472 23. Krol, L. R., Pawlitzki, J., Lotte, F., Gramann, K. & Zander, T. O. Sereega: Simulating event-related eeg activity. *J.*
473 *Neurosci. Methods* **309**, 13–24, [10.1016/j.jneumeth.2018.08.001](https://doi.org/10.1016/j.jneumeth.2018.08.001) (2018).
- 474 24. Pinceti, A., Sankar, L. & Kosut, O. Generation of synthetic multi-resolution time series load data. arXiv:2107.03547
475 (2021).
- 476 25. Yuan, Z., Jiang, Y., An, Z., Ma, W. & Wang, Y. Seismic resolution improving by a sequential convolutional neural network.
477 *PLOS ONE* **19**, e0304981, [10.1371/journal.pone.0304981](https://doi.org/10.1371/journal.pone.0304981) (2024).
- 478 26. Ibarra-Fiallo, J., Lara, J. A. & Agudelo Moreno, D. Cosibd, [10.5281/zenodo.18295713](https://doi.org/10.5281/zenodo.18295713) (2025). Version v2. Dataset.
- 479 27. Welch, P. D. The use of fast fourier transform for the estimation of power spectra: A method based on time averaging
480 over short, modified periodograms. *IEEE Transactions on Audio Electroacoustics* **15**, 70–73, [10.1109/TAU.1967.1161901](https://doi.org/10.1109/TAU.1967.1161901)
481 (1967).

482 Acknowledgments

483 This research was supported by Dean’s Office of the Polytechnic College of the San Francisco de Quito University and partially
484 by ProyExcel-0069 project of the Andalusian University, Research and Innovation Department.

485 Author Contributions

486 J. I. F. handled the methodological design for artificial data creation, probabilistic analysis, spline-based variations, noise
487 distributions, and random node selection. J. A. L. was responsible for the time series methodological design. D. A. M.
488 performed data processing and validation analysis. All of the authors have contributed to writing the manuscript.

489 Competing Interests

490 The authors declare no competing interests.